Supplement S1

Theoretical Background and Derivation of Hypotheses

Range restriction in the predictor X does not bias β_1 or d because removing observations leaves the expectation $E[Y|X]$ – and thus the underlying $X - Y$ relationship – unchanged. In contrast, β_1^{std}, R^2 and r depend on the marginal variance of X, so changing that variance biases theses standardized or variance-based measures (Carretta & Ree, 2022; Sackett et al., 2007; Le et al., 2016). Restricting the outcome Y, however, can bias all estimates. Truncating Y distorts both the error distribution and the joint distribution of X and Y, yielding biased β_1, d and variance-based measures (Bobko et al., 2001; Li, 2015; Carretta & Ree, 2022; Sackett et al., 2007; Le et al., 2016). Consider the multiple linear model

$$Y = \beta X + \varepsilon, \qquad E[\varepsilon \mid X] = 0, \qquad \mathrm{Var}(\varepsilon) = \sigma_\varepsilon^2,$$

with $X = (1, X_1, X_2, X_3, ..., X_p)^t$ including the intercept, the predictor of interest X_1 and additional covariates $(X_2, X_3, ..., X_p)^t$, as well as an exogenous error term ε. Throughout, restriction refers to selection based on the predictor of interest X_1 (or on Y), while other covariates are unrestricted. Let S denote a sample selection indicator,

$$S = \begin{cases} 1 & \text{observation is selected,} \\ 0 & \text{observation is not selected.} \end{cases}$$

The extent to which range restriction affects the estimands depends on how the restriction mechanism is specified and on the resulting reduction in variance. In practice, analysts may operationalize restriction in different ways. Two general approaches can be distinguished: deterministic restriction and probabilistic restriction. Deterministic restriction refers to selection based on fixed cutoffs (e.g., excluding observations below a threshold), whereas probabilistic restriction assigns selection probabilities conditional on the value of X or Y. Importantly, not all restriction mechanisms necessarily reduce variance to the same extent, which may affect the applicability and magnitude of the theoretical results discussed below.

We study two cases:

1. selection depending only on the predictor X_1, i.e. $P(S = 1 \mid X, Y) = P(S = 1 \mid X_1)$

2. selection depending only on the outcome Y, i.e. $P(S = 1 \mid X, Y) = P(S = 1 \mid Y)$.

Unstandardized regression coefficient

The population regression coefficients $\beta = (\beta_0, \beta_1, \beta_2, ..., \beta_p)$ can be written as:

$$\hat{\beta} = (\mathrm{Var}(X))^{-1} \,\mathrm{Cov}(X, Y).$$

In the selected sample ($S = 1$),

$$\hat{\beta}_S = (\mathrm{Var}(X \mid S = 1))^{-1} \,\mathrm{Cov}(X, Y \mid S = 1).$$

Expanding the covariance gives

$$\mathrm{Cov}(X, Y \mid S = 1) = \mathrm{Var}(X \mid S = 1)\, \beta + \mathrm{Cov}(X, \varepsilon \mid S = 1),$$

and therefore

$$\hat{\beta}_S = \beta + (\mathrm{Var}(X \mid S = 1))^{-1} \mathrm{Cov}(X, \varepsilon \mid S = 1).$$

Restriction in the predictor X: If $P(S = 1 \mid X, Y) = P(S = 1 \mid X_1)$, then $E[\varepsilon \mid X, S = 1] = E[\varepsilon \mid X] = 0$, so $\mathrm{Cov}(X, \varepsilon \mid S = 1) = 0$. Thus $\hat{\beta}_S = \beta$, i.e. unstandardized regression coefficients are unchanged by selection on X_1.

Restriction in the outcome Y: If $P(S = 1 \mid X, Y) = P(S = 1 \mid Y)$, conditioning on S changes the distribution of ε: $E[\varepsilon \mid X, S = 1] \neq 0$, and therefore $\mathrm{Cov}(X, \varepsilon \mid S = 1) \neq 0$ in general. Hence $\hat{\beta}_S \neq \beta$ and unstandardized coefficients are biased.

Standardized regression coefficient

The population standardized regression coefficient vector is defined as

$$\beta^{\text{std}} = R_{XX}^{-1} r_{XY},$$

with R_{XX} is the correlation matrix of the predictors `X` and r_{XY} the vector of correlations between `X` and `Y`.

The standardized regression coefficient for X_1 is the first element of this vector:

$$\beta_1^{\text{std}} = (R_{XX}^{-1} r_{XY})_1.$$

The coefficient represents the partial standardized effect of X_1 on `Y`, controlling for $X_2, \ldots, X_p$.

Restriction in the predictor X: Selection on X_1 preserves exogeneity, i.e. $E[\varepsilon \mid X, S = 1] = 0$, but alters the joint distribution of the predictors. In particular, the predictor correlation matrix changes $R_{XX,S} \neq R_{XX}$, and the standardized coefficient becomes $\beta_1^{\text{std},S} = (R_{XX,S}^{-1} r_{XY,S})_1 \neq \beta_1^{std}$.

Restriction in the outcome Y: For restriction in `Y` we have $E[\varepsilon \mid X, S = 1] \neq 0$. Both the predictor–outcome correlations and the covariance structure are distorted. As a result, $\beta^{\text{std},S} = (R_{XX,S}^{-1} r_{XY,S})$ does not correspond to the population standardized effect, and standardized coefficients are biased.

Coefficient of determination

We define R^2 as the population coefficient of determination of the bivariate regression of `Y` on X_1,

$$R^2 = \frac{\beta_1^2 \text{Var}(X_1)}{\text{Var}(Y)}.$$

Restriction in the predictor X: Because restriction on X_1 changes $\text{Var}(X_1)$ and therefore $\text{Var}(Y)$, it follows that $R_S^2 \neq R^2$ in general.

Restriction in the outcome `Y`*:* Selection on `Y` distorts both the variance of `Y` and the error distribution, so the population identity defining R^2 no longer holds and $R_S^2 \neq R^2$.

Intuition: R^2 measures explained variance relative to total variance. Range restriction alters total variance and therefore alters R^2.

Pearson correlation

The Pearson correlation between X_1 and Y is

$$r = \frac{\mathrm{Cov}(X_1, Y)}{\sqrt{\mathrm{Var}(X_1)\mathrm{Var}(Y)}} = \frac{\beta_1\sqrt{\mathrm{Var}(X_1)}}{\sqrt{\beta_1^2\mathrm{Var}(X_1) + \mathrm{Var}(\varepsilon)}}.$$

Restriction in the predictor X_1*:* Because r depends on $\mathrm{Var}(X_1)$ and $\mathrm{Var}(Y)$, restriction on X_1 generally implies $r_S \neq r$.

Restriction in the outcome Y*:* Restriction on Y alters both $\mathrm{Var}(Y)$ and the error distribution, so $r_S \neq r$ in general.

Cohen's d

Let $X_1 \in \{0, 1\}$ be a binary predictor and assume $Y \mid X_1 = x_1 \sim \mathcal{N}\mu_{x_1}, \sigma^2)$. The population standardized mean difference is

$$\delta = \frac{\mu_1 - \mu_0}{\sigma}.$$

Cohen's d is the sample estimator

$$\hat{d} = \frac{\bar{Y}_1 - \bar{Y}_0}{s_p}, \qquad s_p^2 = \frac{(n_0 - 1)s_0^2 + (n_1 - 1)s_1^2}{n_0 + n_1 - 2}.$$

Restriction in the predictor X: If selection depends only on X_1, then within each group $Y \mid X_1 = x_1, S = 1 \overset{d}{=} Y \mid X_1 = x_1$. Thus the group means μ_{x_1} and common variance σ^2 are unchanged, and the population effect size satisfies $\delta_S = \delta$. Conditional on the selected group sizes, $\hat{d}_S$ has the same sampling distribution as $\hat{d}$ from an unrestricted sample of the same sizes.

Restriction in the outcome Y: If selection depends on Y (e.g. truncation $Y > c$), then

$Y \mid X_1 = x_1, S = 1$ is a truncated normal distribution. Its mean and variance are

$$E[Y \mid X_1 = x_1, S = 1] = \mu_{x_1} + \sigma\lambda\left(\frac{c - \mu_{x_1}}{\sigma}\right),$$

$$\mathrm{Var}(Y \mid X_1 = x_1, S = 1) = \sigma^2\left\{1 - \kappa\left(\frac{c - \mu_{x_1}}{\sigma}\right)\right\},$$

with the usual inverse Mills ratio λ and variance correction κ. Both depend on μ_{x_1}, so the mean difference and group variances change. Hence $\delta_S \neq \delta$ and Cohen's d estimates a different population effect size after selection on `Y`.

Derivation of true values

For β_1 the true value is known from the DGP and was therefore set directly to this value.
β_1^{std} was computed as

$$\beta_1^{std} = (R_{XX}^{-1} r_{XY})_1,$$

where R_{XX} denotes the population correlation matrix of the predictors and r_{XY} the vector of correlations between each predictor and the outcome. The subscript $(\cdot)_1$ refers to the element corresponding to the focal predictor.
The Pearson correlation coefficient r was computed as

$$r = \beta_1 \frac{\sqrt{Var(X)}}{\sqrt{Var(Y)}},$$

and the coefficient of determination R^2 was computed as

$$R^2 = \beta_1^2 \frac{Var(X)}{Var(Y)}$$

$Var(X)$, $Var(Y)$, R_{XX} and r_{XY} were obtained as Monte-Carlo averages of the empirical variances computed from 1,000 independent simulated datasets generated under the unrestricted data-generating process.
If `X` was binary, Cohen's d was computed as

$$d = \frac{\mu_1 - \mu_0}{s_p},$$

where μ_1 and μ_0 denote the population means of `Y` in the two groups and s_p is the pooled standard deviation of `Y` across groups. These quantities were approximated by Monte Carlo averages across 1,000 independent simulated datasets.

Effect Sizes

The core estimation procedure was an ordinary least squares (OLS) regression of `Y` on `X` (and any additional covariates specified in the DGP), implemented via `stats::lm()` in R with default settings. Depending on the type of predictor `X`, the following effect sizes were computed:

For **numeric `X`**, a linear regression model

$$Y \sim X \text{ (+ additional covariates)}$$

was fitted using `lm()` to derive the estimated unstandardized regression coefficient for β_1. In addition, a standardized version of this model was fitted after `Y` and the predictors were z-standardized within the script to derive an estimate for the standardized regression coefficient β_1^{std}. Furthermore, an estimate for the partial coefficient of determination R^2 was computed as:

$$\hat{R^2} = \frac{\hat{\beta_1}^2 \widehat{Var(X)}}{\widehat{Var(Y)}}$$

with $\hat{\beta_1}$ denotes the slope estimate from the fitted linear model and $\widehat{Var(X)}$ and $\widehat{Var(Y)}$ denotes the empirical variances computed from the simulated dataset. In addition, the Pearson correlation coefficient r was estimated as:

$$\hat{r} = \frac{\widehat{Cov(X,Y)}}{\sqrt{\widehat{Var(X)}}\sqrt{\widehat{Var(Y)}}} = \frac{\hat{\beta_1}\sqrt{\widehat{Var(X)}}}{\sqrt{\widehat{Var(Y)}}}$$

implying that $r = sign(\hat{\beta_1})\sqrt{\hat{R^2}}$.

Note: In the univariate case, r and R^2 coincide with the standardized regression coefficient. In multivariate models, however, the r does not represent a partial effect, as it does not adjust for the presence of additional predictors.

For **binary `X`** (two groups), the same linear model and formulas was used to derive β_1, β_1^{std} and

R^2. Instead of r the Cohen's d was computed as:

$$\hat{d} = \frac{\hat{\bar{Y}}_1 - \hat{\bar{Y}}_0}{\hat{s}}$$

with

$$\hat{s} = \sqrt{\frac{(n_0 - 1)\hat{s_0}^2 + (n_1 - 1)\hat{s_1}^2}{n_0 + n_1 - 2}}),$$

with $\hat{\bar{Y}}_0$ and $\hat{\bar{Y}}_1$ denoting the group means, n_0 and n_1 group sizes, and $\hat{s_0}^2$ and $\hat{s_1}^2$ group variances in the simulated datasets.

Performance measures

Let θ denote the true value of a given effect size measure and $\hat{\theta}_i$ the corresponding estimate in simulation i, for $i = 1, \ldots, n_{\text{sim}}$.

Absolute and relative bias. The absolute bias of the estimator was computed as:

$$\widehat{\text{AbsBias}} = \frac{1}{n_{\text{sim}}} \sum_{i=1}^{n_{\text{sim}}} \left(\hat{\theta}_i - \theta\right).$$

For $\theta \neq 0$, we defined the relative bias in simulation i as

$$b_i = \frac{\hat{\theta}_i - \theta}{\theta},$$

and summarized it across replications by the mean relative bias

$$\widehat{\text{RelBias}} = \frac{1}{n_{\text{sim}}} \sum_{i=1}^{n_{\text{sim}}} b_i.$$

The Monte Carlo standard error (MCSE) of the mean relative bias was computed as

$$\text{MCSE}_{\widehat{\text{RelBias}}} = \frac{S_b}{\sqrt{n_{\text{sim}}}},$$

where S_b is the sample standard deviation of $(b_1, \ldots, b_{n_{\text{sim}}})$. The MCSE for the absolute bias was calculated analogously by replacing *RelBias* with *AbsBias* and using the corresponding sample standard deviation.

Coverage, power, and power loss. For each effect size measure θ, except R^2, and each simulation i, we computed a two-sided 95% confidence interval $[L_i, U_i]$ for the parameter of interest. Coverage was assessed by the indicator

$$C_i = \mathbb{I}\{\theta \in [L_i, U_i]\},$$

and the empirical coverage probability was estimated by

$$\widehat{\text{Cov}} = \frac{1}{n_{\text{sim}}} \sum_{i=1}^{n_{\text{sim}}} C_i.$$

Using the same confidence intervals, we defined rejection of the null hypothesis of no effect by

$$R_i = \mathbb{I}\{0 \notin [L_i, U_i]\}.$$

In conditions where the true effect is zero, the mean of R_i estimates the type I error rate; in conditions with a non-zero true effect, it estimates power. We computed these rejection rates both for the convenience samples and for matched-size random samples.

Aggregated measures

For the pre-specified subjective effect-size categories (weak, moderate, strong), the unweighted mean and the empirical of the DGP-specific mean relative biases was computed as:

$$\bar{b}_{r,\beta} = \frac{1}{D} \sum_{d=1}^{D} \bar{b}_{d,r,\beta},$$

where D denotes the number of valid DGPs contributing to the condition.

Uncertainty of the aggregated mean bias was quantified using the 95% confidence interval based on the standard errors across DGPs:

$$CI_{0.95} = \bar{b}_{r,\beta} \pm 1.96 \frac{s_{r,\beta}}{\sqrt{D}}$$

where $s_{r,\beta}$ denotes the standard deviation of DGP-specific mean biases within the respective condition.

Partial R^2 and variance reduction

An estimate of the partial coefficient of determination R^2 was computed as

$$\widehat{R^2} = \frac{\hat{\beta}_1^2 \widehat{\mathrm{Var}}(X)}{\widehat{\mathrm{Var}}(Y)}.$$

Here, $\hat{\beta}_1$ denotes the slope estimate from the fitted linear model, and $\widehat{\mathrm{Var}}(X)$ and $\widehat{\mathrm{Var}}(Y)$ denote the empirical variances computed from the simulated data.

Variance reduction was estimated as

$$\mathrm{Reduction} = 1 - \frac{\mathrm{Var}_{\mathrm{restricted}}}{\mathrm{Var}_{\mathrm{unrestricted}}}.$$

Both estimates are based on 1,000 Monte Carlo simulations.

Workflow, Implementation, and Inclusion Procedures

Analysts workflow

Analysts received an instruction document, an R or R Markdown script template, a helper script, and ADEMP templates for both the confirmatory part and the overall many-analysts study. The full analyst instructions, templates, and helper scripts are available on the OSF.

The confirmatory component was designed to allow controlled RDF while preserving comparability across analysts. Analysts were asked to implement a data-generating process (DGP) in a predefined function, `dgp(beta1, restrict, n)`, and to specify a corresponding analysis model formula. Within this structure, analysts could choose the distribution of the predictor `X`, include additional covariates, define weak, moderate, and strong effect-size scenarios, choose the sample size, and implement weak, moderate, and strong restriction mechanisms for `X` and `Y`. At the same time, the confirmatory component required the true data-generating model to be a correctly specified linear model.

The exploratory component allowed analysts to investigate additional design choices or violations of the confirmatory assumptions. Analysts were invited to extend the simulations, for example by introducing measurement error, heteroscedasticity, nonlinearity, model misspecification, additional effect-size measures, or alternative performance metrics. They were asked to document these exploratory choices in a short written report.

Reproducibility procedures

To support reproducibility and comparability, analysts were provided with pre-written simulation and analysis code. The helper script implemented the main simulation workflow, computed effect-size estimates, bias measures, power-related quantities, and produced standardized output figures. Analysts were instructed to modify only the designated parts of the confirmatory template and to submit their completed script together with the resulting figures and reproducibility information.

The helper script set a fixed random seed and used standardized functions for validating the DGP, running the simulations, deriving true values, computing performance measures, and visualizing results. Analysts were asked to run the simulation for 10,000 Monte Carlo iterations after testing their implementation with a smaller number of iterations. For the aggregated analyses reported in the main manuscript, all valid submitted DGPs were re-run centrally using a common computational setup.

Analysts were also asked to submit a `sessionInfo()` output file documenting the R version and package versions used for their analysis. In addition, the lead team used a version-controlled R environment managed with `renv`; the corresponding lockfile is publicly archived.

Code review and inclusion criteria

All submitted scripts were independently reviewed by two reviewers. The reviewers checked whether analysts met the eligibility criteria, whether the submitted code followed the confirmatory instructions, whether the DGP and model formula were correctly specified, and whether the code ran successfully with the provided helper functions.

Analysts were eligible to participate if they had at least Master's-level training and prior experience with simulation-based research and R. A submitted confirmatory analysis was eligible for inclusion if the analyst followed the required structure of the confirmatory component, implemented a valid `dgp()` function, used the required restriction levels, provided non-zero values for the three effect-size scenarios, specified a valid analysis model, and submitted reproducible code.

Submissions were excluded from the confirmatory aggregation if the code could not be executed, if the required output structure was not provided, if the DGP violated the confirmatory design constraints, or if the quality-control checks indicated that the implementation was not valid. Disagreements between reviewers were resolved by discussion.

Quality-control checks

The helper script implemented several automated quality-control checks before running the full simulation. These checks verified that `beta1` was a numeric vector of length three with non-zero entries, that `restrict` contained exactly the four levels `none`, `weak`, `moderate`, and `strong`, and that the model formula used `y` as the outcome and included `x` as a predictor.

The output of the dgp() function was checked for the required list structure, containing exactly two data frames named convenience_x and convenience_y. Each data frame had to contain at least the variables x and y, with y numeric and x either numeric or binary. The checks further required sufficient variation in x and y, a sufficient number of unique values for numeric variables, equal sample sizes for the samples restricted on X and Y, and identical output for convenience_x and convenience_y when restrict = "none".

In addition, the helper script fitted the specified analysis model to generated datasets and checked whether the focal coefficient could be estimated. As a benchmark, it also evaluated whether random samples generated under the unrestricted condition produced approximately unbiased estimates of the unstandardized regression coefficient at the chosen sample size. These checks were intended to detect implementation errors and DGP specifications that violated the assumptions of the confirmatory component before the full simulation was run.

Computational environment

The helper script required the R packages ggplot2 and dplyr. Missing packages were installed automatically by the helper script, and all required packages were loaded before running the simulations. Analysts were instructed to submit their completed script, output figures, and a sessionInfo() file.

For the aggregated analyses, the lead team re-ran all included DGPs in a version-controlled R environment managed using renv. This ensured that the final aggregation was performed under a consistent computational environment. The corresponding renv lockfile and analysis scripts are archived together with the study materials.

Meta regression

For each effect size measure and convenience-sample type, the outcome in the meta-regression models was the DGP-level mean relative bias. Let y_i denote the mean relative bias for observation i, where each observation corresponds to one condition within one analyst-specified DGP. Because MCSEs were negligible with 10,000 replications, weighting by simulation uncertainty would have had minimal impact on the results. More importantly, MCSEs depend not only on the number of replications but also on the empirical variability of the estimator. As a result, weighting schemes based on MCSE would give greater influence to DGPs with lower empirical variability or more replications. Since our goal was to characterize variability across DGPs rather than simulation precision, we therefore assigned equal weight to all observations. Accordingly, we fitted unweighted multilevel meta-regression models of the form:

$$y_i = \beta_0 + \sum_{k=1}^{K} \beta_k x_{ki} + u_{d_{(i)}} + v_{\alpha_{(i)}} + e_i$$

where x_{ki} denotes the fixed-effect covariates, $u_{d_{(i)}}$ the random intercept for DPG, $v_{\alpha_{(i)}}$ the random intercept for analyst, and e_i the sampling error term. We assume:

$$u_{d_{(i)}} \sim N(0, \tau^2_{DGP}), \qquad v_{\alpha_{(i)}} \sim N(0, \tau^2_{analyst}), \qquad e_i \sim N(0, V_i), \tag{1}$$

with $V_i = 1$ for all observations. Thus, each DGP contributed equally to the estimation. The analyst-level random intercept was included only if at least one analyst contributes more than one DGP, for example by submitting separate designs for numeric and binary predictors. Otherwise the model reduced to:

$$y_i = \beta_0 + \sum_{k=1}^{K} \beta_k x_{ki} + u_{d_{(i)}} + e_i.$$

For the unstandardized regression β_1, the standardized regression coefficient β_1^{std} and Cohen's d, we fitted random-effects meta-regression models using `metafor::rma.mv`. Two families of moderator models were estimated.

First, the **RDF-related models** examined whether variation in relative bias is associated with researcher degrees of freedom in the simulation design. The subjective RDF model includes analysts' categorical design choices as predictors:

$$
\begin{aligned}
y_i = \beta_0 &+ \beta_1 \,\mathrm{RestrictionLevel}_i + \beta_2 \,\mathrm{EffectSizeStrength}_i \\
&+ \beta_3 \,\mathrm{ConvenienceSampleSize}_i + \beta_4 \,\mathrm{PredictorType}_i + \beta_5 \,\mathrm{DistributionType}_i \\
&+ \beta_6 \,\mathrm{RestrictionMechanism}_i + \beta_7 \,\mathrm{CovariatesIncluded}_i + \beta_8 \,\mathrm{CorrelatedCovariates}_i \\
&+ u_{d(i)} + v_{a(i)} + e_i .
\end{aligned}
$$

Here, *RestrictionLevel* refers to the subjective restriction categories weak, moderate, and strong; *EffectSizeStrength* refers to the subjective effect-size categories weak, moderate, and strong; and the remaining predictors indicate convenience-sample size, predictor type (numeric vs. binary), distribution type (parametric vs. non-parametric), restriction mechanism (deterministic vs. probabilistic), inclusion of covariates, and inclusion of correlated covariates. For Cohen's d, predictor type was omitted because this effect size is only defined for binary predictors.

The objective RDF model replaces the subjective design labels by quantitative characteristics of the implemented DGP:

$$
\begin{aligned}
y_i = \beta_0 &+ \beta_1 \,\mathrm{VarianceReduction}_i + \beta_2 \,\mathrm{TruePartialR}^2_i \\
&+ \beta_3 \,\mathrm{ConvenienceSampleSize}_i + \beta_4 \,\mathrm{PredictorType}_i + \beta_5 \,\mathrm{DistributionType}_i \\
&+ \beta_7 \,\mathrm{CovariatesIncluded}_i + \beta_8 \,\mathrm{CorrelatedCovariates}_i + u_{d(i)} + v_{a(i)} + e_i .
\end{aligned}
$$

Again, predictor type was omitted for Cohen's d. In both RDF models, categorical predictors were dummy-coded with appropriate reference categories.

Second, the **prior-belief model** examined whether analysts' pre-simulation plausibility ratings are associated with relative bias. Let $\mathrm{Plausibility}_i$ denote the relevant plausibility rating for the hypothesis under consideration. The model is:

$$y_i = \beta_0 + \beta_1 \cdot \text{Plausibility}_i + u_{d_{(i)}} + v_{\alpha_{(i)}} + e_i,$$

or, if no analyst contributes more than one DGP,

$$y_i = \beta_0 + \beta_1 \cdot \text{Plausibility}_i + u_{d_{(i)}} + e_i.$$

Categories with insufficient observations, such as rarely selected RDF settings or plausibility levels, were excluded from the corresponding models to ensure estimability. Because R^2 and Pearson's r are redundant with the standardized regression coefficient, no separate meta-regressions were estimated for these measures.

References

Bobko, P., Roth, P. L., & Bobko, C. (2001). Correcting the effect size of d for range restriction and unreliability. *Organizational Research Methods*, *4*(1), 46-61. doi: https://doi.org/10.1177/109442810141003

Carretta, T. R., & Ree, M. J. (2022). Correction for range restriction: Lessons from 20 research scenarios. *Military Psychology*, *34*(5), 551–569. doi: https://doi.org/10.1080/08995605.2021.2022067

Le, H., Oh, I.-S., Schmidt, F. L., & Wooldridge, C. D. (2016). Correction for range restriction in meta-analysis revisited: Improvements and implications for organizational research. *Personnel Psychology*, *69*(4), 975–1008. doi: https://doi.org/10.1111/peps.12122

Li, J. C.-H. (2015). Cohen's d corrected for case iv range restriction: A more accurate procedure for evaluating subgroup differences in organizational research. *Personnel Psychology*, *68*(4), 899-927. Retrieved from `https://onlinelibrary.wiley.com/doi/abs/10.1111/peps.12096` doi: https://doi.org/10.1111/peps.12096

Sackett, P. R., Lievens, F., Berry, C. M., & Landers, R. N. (2007). A cautionary note on the effects of range restriction on predictor intercorrelations. *Journal of Applied Psychology*, *92*(2), 538–544. doi: https://doi.org/10.1037/0021-9010.92.2.538